

Developing Kernel Drivers With Modern C++

Pavel Yosifovich

@zodiacon



zodiacon@live.com



About Me

2

- Developer, Trainer, Author and Speaker
- Book author
 - “Windows Kernel Programming” (2019)
 - “Windows Internals 7th edition, Part 1” (co-author, 2017)
 - “WPF 4.5 Cookbook” (2012)
- Pluralsight author
- Author of several open-source tools (<http://github.com/zodiacon>)
- Blogs: <http://blogs.microsoft.co.il/pavely>, <http://scorpiosoftware.net>



Agenda

3

- Kernel Programming
- User Mode vs. Kernel Mode
- C++ in the Kernel
- Summary
- Q & A

Kernel Programming

4

- Kernel programming refers to Kernel device drivers running in kernel mode (and kernel space)
- Classically written in C
- All Microsoft driver samples are still written in C
- Since 2012, Microsoft officially supports using C++ in drivers
 - Shame MS not taking advantage of this themselves (at least not publicly)
- However, using C++ in kernel mode is not the same as user mode

User Mode vs. Kernel Mode Development


5

	User Mode	Kernel Mode
→ Unhandled Exceptions	Unhandled exceptions crash the process	Unhandled exceptions crash the system
→ Termination	When process terminates, all private memory and resources are freed automatically	If driver unloads without freeing everything it was using, there is a leak, only resolved in the next boot
→ Return values	API errors sometimes ignored	Should (almost) never ignore errors
→ IRQL	Always <code>PASSIVE_LEVEL</code> (0)	May be <code>DISPATCH_LEVEL</code> (2) or higher
→ Bad coding	Typically localized to the process	Can have system-wide effect
→ Testing and Debugging	Typically testing and debugging done on the developer machine	Must use another machine for testing and debugging
→ Libraries	Can use almost any C/C++ library (e.g. STL)	Most standard libraries cannot be used
→ Exception Handling	Can use C++ exceptions or Structured Exception Handling (SEH)	Only SEH can be used
→ C++ Usage	Full C++ runtime available	No C++ runtime

C++ in the Kernel

6

- No runtime
 - No exception handling
 - No new/delete operators
 - No constructors called for global objects
 - The standard library is (almost) useless
- All other features work normally
 - auto, range-based for, lambdas, ...
 - Constructors / destructors, virtual functions
 - Templates of all sorts
 - Move semantics
 - Operator overloading



RAII

7

- Resource Acquisition Is Initialization
- C++ cleanup pattern
- Just as important (if not more so) in the kernel

RAII Example

8

```
template<typename TLock>
struct AutoLock {
    AutoLock(TLock& lock) : _lock(lock) {
        _lock.Lock();
    }

    ~AutoLock() {
        _lock.Unlock();
    }

private:
    TLock& _lock;
};
```

```
class FastMutex {
public:
    void Init();

    void Lock();
    void Unlock();

private:
    FAST_MUTEX _mutex;
};
```

```
AutoLock<FastMutex> locker(MyMutex); // C++ 14
```

```
AutoLock locker(MyMutex); // C++ 17
```

```
void FastMutex::Init() {
    ExInitializeFastMutex(&_mutex);
}

void FastMutex::Lock() {
    ExAcquireFastMutex(&_mutex);
}

void FastMutex::Unlock() {
    ExReleaseFastMutex(&_mutex);
}
```


More RAI Examples

9

- Handles
- (in Mini-filter) File Name Information
- (in Mini-filter) Contexts
- Memory allocations

Dynamic Memory Allocations

10

- Drivers allocate memory from the kernel memory pools
 - Non paged pool or paged pool
- Use **ExAllocatePoolWithTag** to allocate
 - Tag is a 4-byte value useful for identifying the allocating driver (or a component that is part of a driver)
 - Typically specified as a 4-character ASCII string
 - Can view with the *Poolmon* WDK tool or the kernel debugger
- Use **ExFreePool** to free the buffer
- How to invoke ctor/dtor?

```
PVOID ExAllocatePoolWithTag (  
    _In_ POOL_TYPE PoolType,  
    _In_ SIZE_T NumberOfBytes,  
    _In_ ULONG Tag);
```

Memory Allocation Example (1)

11

```
void* operator new(size_t size, POOL_TYPE type, ULONG tag = 0);
```

```
void* operator new(size_t size, POOL_TYPE type, ULONG tag) {  
    auto p = tag == 0 ? ExAllocatePool(type, size) : ExAllocatePoolWithTag(type, size, tag);  
    if (p == nullptr) {  
        KdPrint(("Failed to allocate %d bytes\n", size));  
    }  
    return p;  
}
```

```
template<typename T>  
using vector = kvector<T, DRIVER_TAG>;  
  
struct MyData {  
    int Count;  
    FastMutex Lock;  
    vector<PDEVICE_OBJECT> Devices;  
};
```

```
auto data = new (NonPagedPool, DRIVER_TAG) MyData;
```


Memory Allocation Example (2)

12

```
// placement new  
void* operator new(size_t size, void* p);
```

```
void* operator new(size_t, void* p) {  
    return p;  
}
```

```
auto data = static_cast<MyData*>(ExAllocatePoolWithTag(  
    NonPagedPool, sizeof(MyData), DRIVER_TAG));  
new (data) MyData;
```

```
data->~MyData();  
ExFreePool(data);
```

Strings

13

- Most kernel APIs work with UNICODE_STRING structures
 - Length and MaximumLength are in bytes
- UNICODE_STRING is not necessary the owner of its characters
 - An owner-focused wrapper will not be useful
- Can create a generic string class
- Provide conversion to a UNICODE_STRING
 - Maintain ownership

```
typedef struct _UNICODE_STRING {  
    USHORT Length;  
    USHORT MaximumLength;  
    PWCH Buffer;  
} UNICODE_STRING;  
typedef UNICODE_STRING *PUNICODE_STRING;  
typedef const UNICODE_STRING *PCUNICODE_STRING;
```

Strings

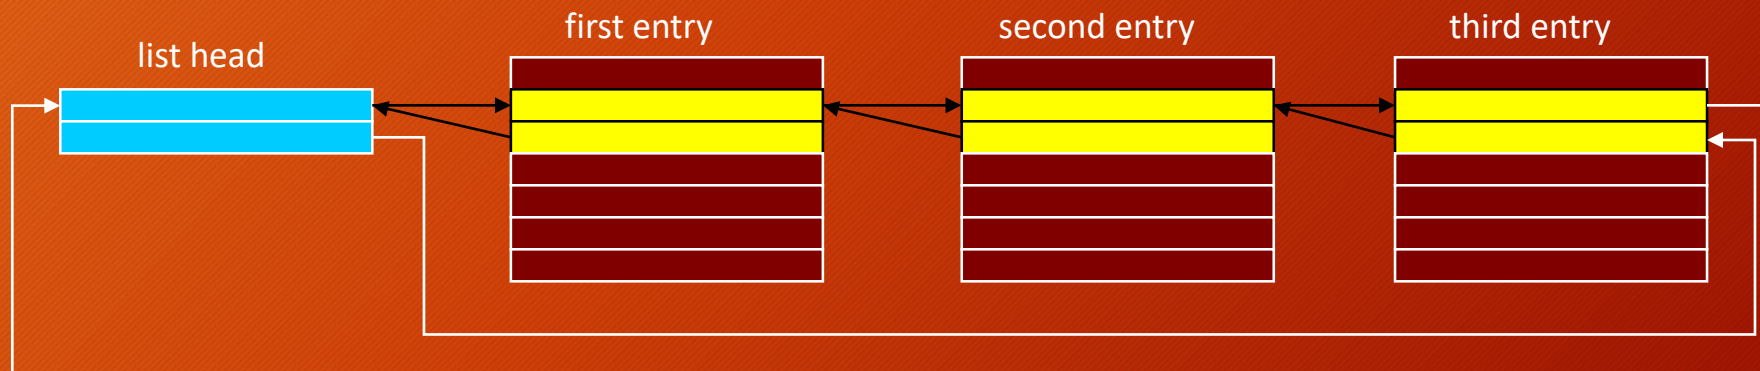
14

Demo

Lists and Queues

15

- The kernel uses circular doubly-linked lists in many data structures
 - LIST_ENTRY structure
 - API is fully documented
 - Drivers can use the same structures
- Normally part of a larger structure
- WDK headers provide the CONTAINING_RECORD macro



Linked Lists

16

```
#include "AutoLock.h"
#include "FastMutex.h"

template<typename T, typename TLock = FastMutex>
class LinkedList {
public:
```

```
    struct MyData {
        int Data;
        LIST_ENTRY Entry; // must be named "Entry"
        int MoreData;
    };
private:
```

```
    LinkedList<MyData> MyList;
```

```
};

MyList.Init();
auto item = new (PagedPool, DRIVER_TAG) MyData;
item->Data = 8;
item->MoreData = 42;

MyList.PushBack(item);
```

```
// expects a LIST_ENTRY named "Entry"
```

```
void PushBack(T* item) {
    AutoLock locker(_lock);
    InsertTailList(&_amp;_head, &item->Entry);
}
```

```
void PushFront(T* value) {
    AutoLock locker(_lock);
    InsertHeadList(&_amp;_head, &item->Entry);
}
```

```
T* RemoveHead() {
    AutoLock locker(_lock);
    auto entry = RemoveHeadList(&_amp;_head);
    return CONTAINING_RECORD(entry, T, Entry);
}
```

```
T* GetHeadItem() {
    AutoLock locker(_lock);
    auto entry = _amp;_head->Flink;
    return CONTAINING_RECORD(entry, T, Entry);
}
```

Components

17

- Some scenarios require objects to be
 - Reference counted
 - Support functionality querying
- The Component Object Model (COM) provides these ideas
 - Implementation in user mode (ATL, WRL)
- What about something like `std::shared_ptr<>`?
 - Not good enough
 - No way to query for functionality
- Solution: implement it ourselves!

Kernel COM

18

```
struct IFilter abstract : IComponent {
    enum { IID = 10 };
};

struct MyFilter : ComponentBase<IFilter, IProcessNotify> {
    template<POOL_TYPE type = PagedPool, ULONG tag = 0>
    static IComponent* Create() {
        return static_cast<IFilter*>(new (type, tag) MyFilter);
    }
};

void FilterRequest(PIRP_Trap) override {
    // do something
}

void ProcessCreated(HANDLE) override {
    // do something
}

private:
    MyFilter() {}
};

IFilter* filter;

auto c = MyFilter::Create<NonPagedPool>();
if (c == nullptr)
    return STATUS_INSUFFICIENT_RESOURCES;

c->QueryInterface(IFilter::IID, reinterpret_cast<void**>(&filter));
if (filter) {
    filter->FilterRequest(...);
    filter->Release();
}
```

Summary

19

- Kernel code no longer stuck with C
- C++ fully supported
- Using C++ properly in the kernel can lead to more robust drivers
- There is no going back!

Thank you!

